DTU Compute
Department of Applied Mathematics and Computer Science

High-Performance Computing

DTU Unix Systems & High-Performance Computing

Using the DTU Unix systems
for your computations

High-Performance Computing

Bernd Dammann

Assoc. Professor, Scientific Computing
DTU Compute, building 303B

MSc in Physics, PhD in Phys. Chemistry
work with DTU's central HPC (and G-bar) since 2001
co-founder of DTU Compute GPUlab (2007)
teach HPC courses (02614, 02616, 41391) and related topics

High-Performance Computing

DTU Computing Center (DCC)

Bernd Dammann

HPC Architect, Consultant & Scientific Lead
DTU Computing Center, building 324

DCC started in 2008
in-sourcing of operations of DTU's central HPC (and G-bar)
HPC consulting for DTU users
HPC Competence Center (since 2013)
GPUlab activities (since 2013)

DTU Computing Center (DCC)

What are we?

- DCC is a (central) DTU unit
- we are **not** part of DTU Compute
  - ... but we have our office at DTU Compute(!)
- we provide ‘free’ access to HPC resources
- we offer HPC hosting services for departments at DTU (e.g. Compute, Chemistry, Environment, ...)
- we offer HPC consulting services for all of DTU
  - ... free of charge (up to a certain extent)
  - ... as part of larger projects

DTU Computing Center (DCC)

Who are we?

- Andrea: HPC expert
- Bernd: HPC architect, expert, scient. lead
- Hans Henrik: HPC expert (GPUs)
- Henning: team head
- Ian: system admin
- Jette: contract administrator
- Nick: HPC expert
- Pietro: project coordinator
- Sebastian: tech lead & system admin

High-Performance Computing

DTU Computing Center (DCC)

Our mission: “make HPC work”

- *easy* access to HPC resources
- provide the support needed – on all levels
  - first steps for newbies
  - performance analysis and tuning
  - new architectures: testing, consulting & porting
- guide users in the HPC landscape
  - local, national and international
- if you use our services, please mention it in your publications (e-mail support@hpc.dtu.dk for details)

DTU Computing Center

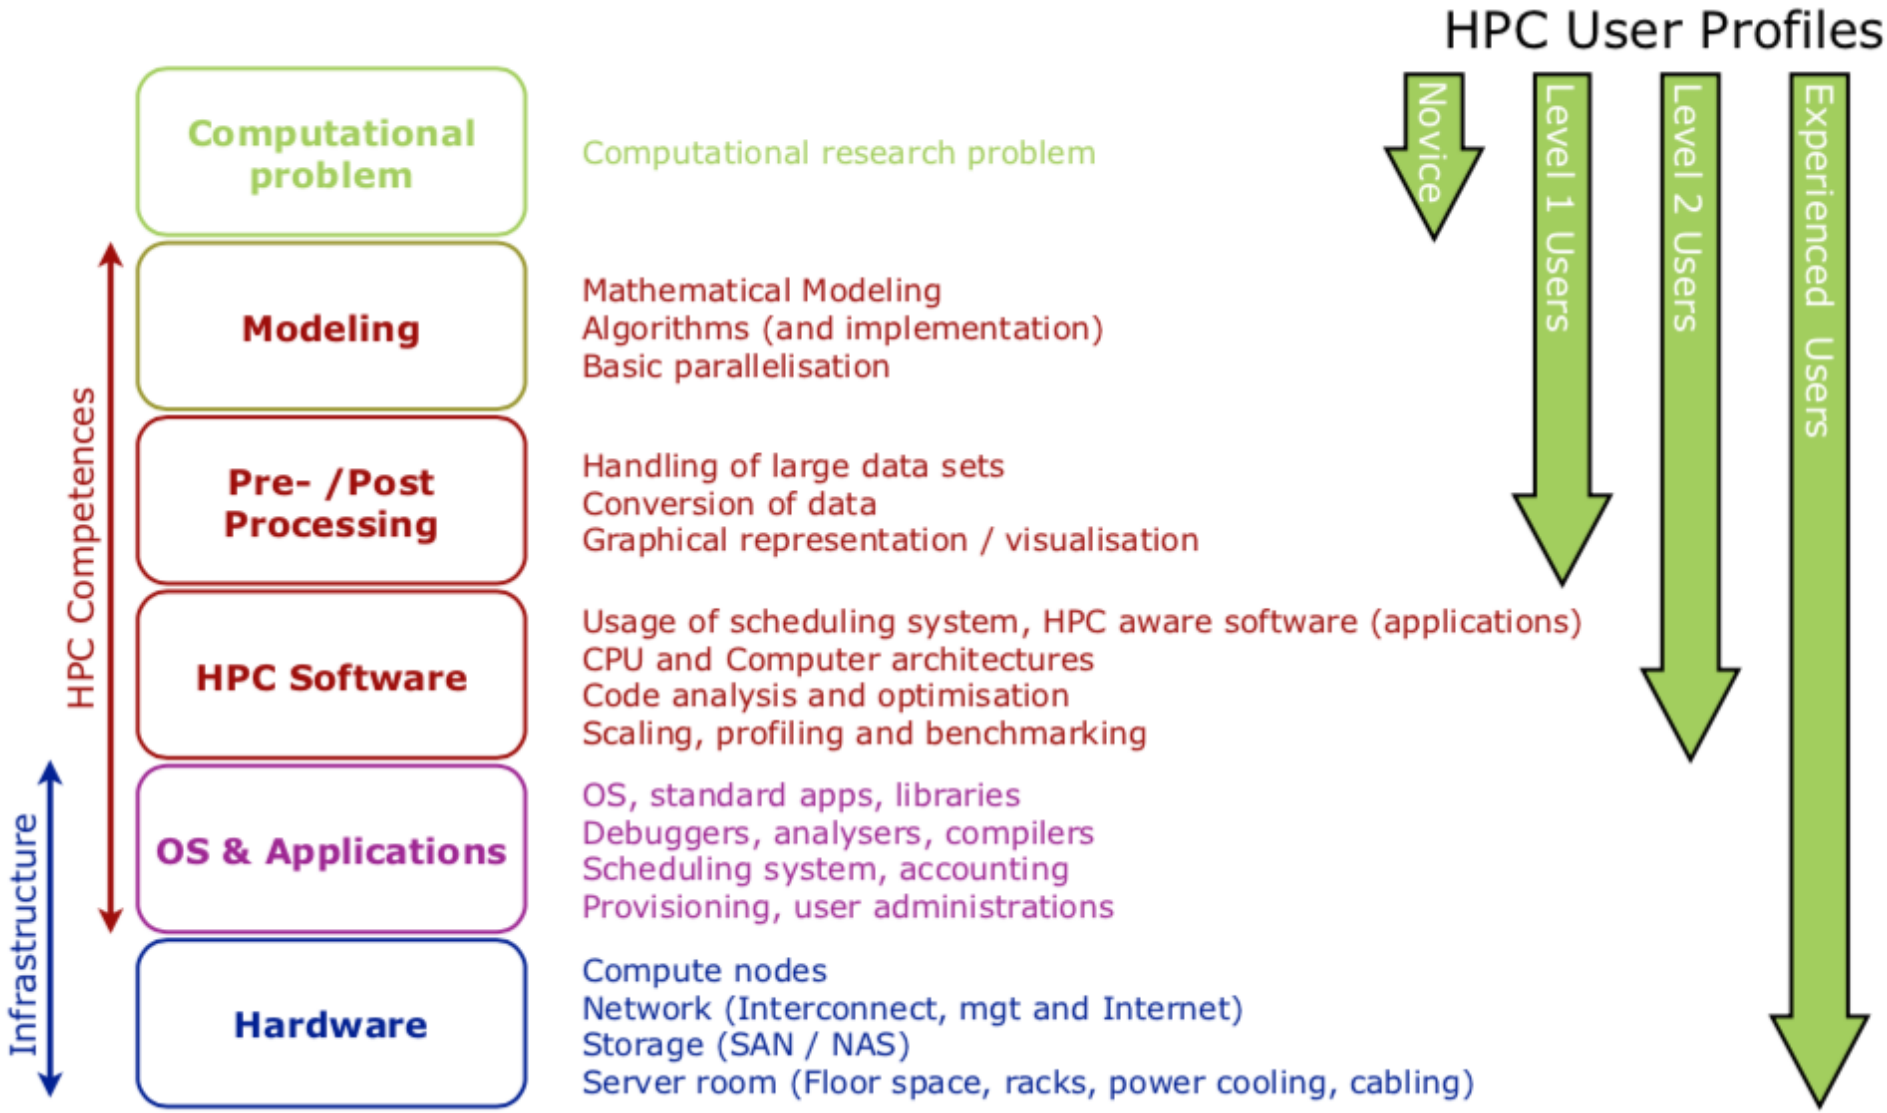

Today's goal

- There is a world 'outside' your laptop – so let's make the first steps!


YOU ARE LEAVING THE
COMFORT ZONE
вы выезжаете из зона
комфорта
VOUS SORTEZ DU ZONE DE
CONFORT
SIE VERLASSEN DEN WOHLFÜHLBEREICH

The DTU Unix systems

a remote databar

also known as G-bar

The DTU Unix systems

Access to the system

- Remote access – from everywhere:
  - ThinLinc remote desktop session:
    - download ThinLinc client from www.thinlinc.com
    - connect to thinlinc.gbar.dtu.dk
    - preferred way, if you work a lot with GUIs
    - on mobile devices: https://thinlinc.gbar.dtu.dk/
  - Secure SHell (ssh) connection (login2.gbar.dtu.dk)
    - for the command line oriented users :-)
- On Campus:
  - ThinLinc from Windows computers

The DTU Unix systems

- But why – and when – should you use it?
- Sooner or later you ...
  - want to solve problems that don't fit your laptop any longer!
  - have to solve many problems, and you run out of time!
  - want to use a software, that doesn't run on your PC.
- So why not start today ... ???
  - it's free!
  - trust me: *"it doesn't hurt!"*

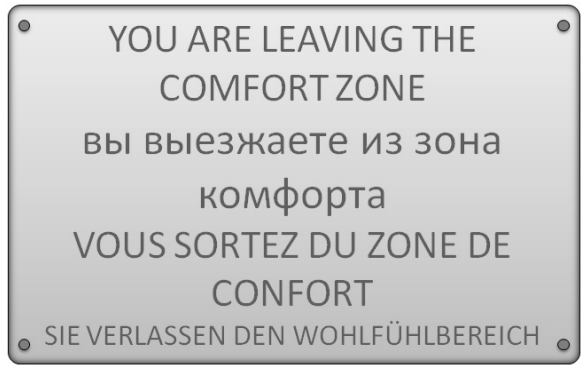

Our Setup

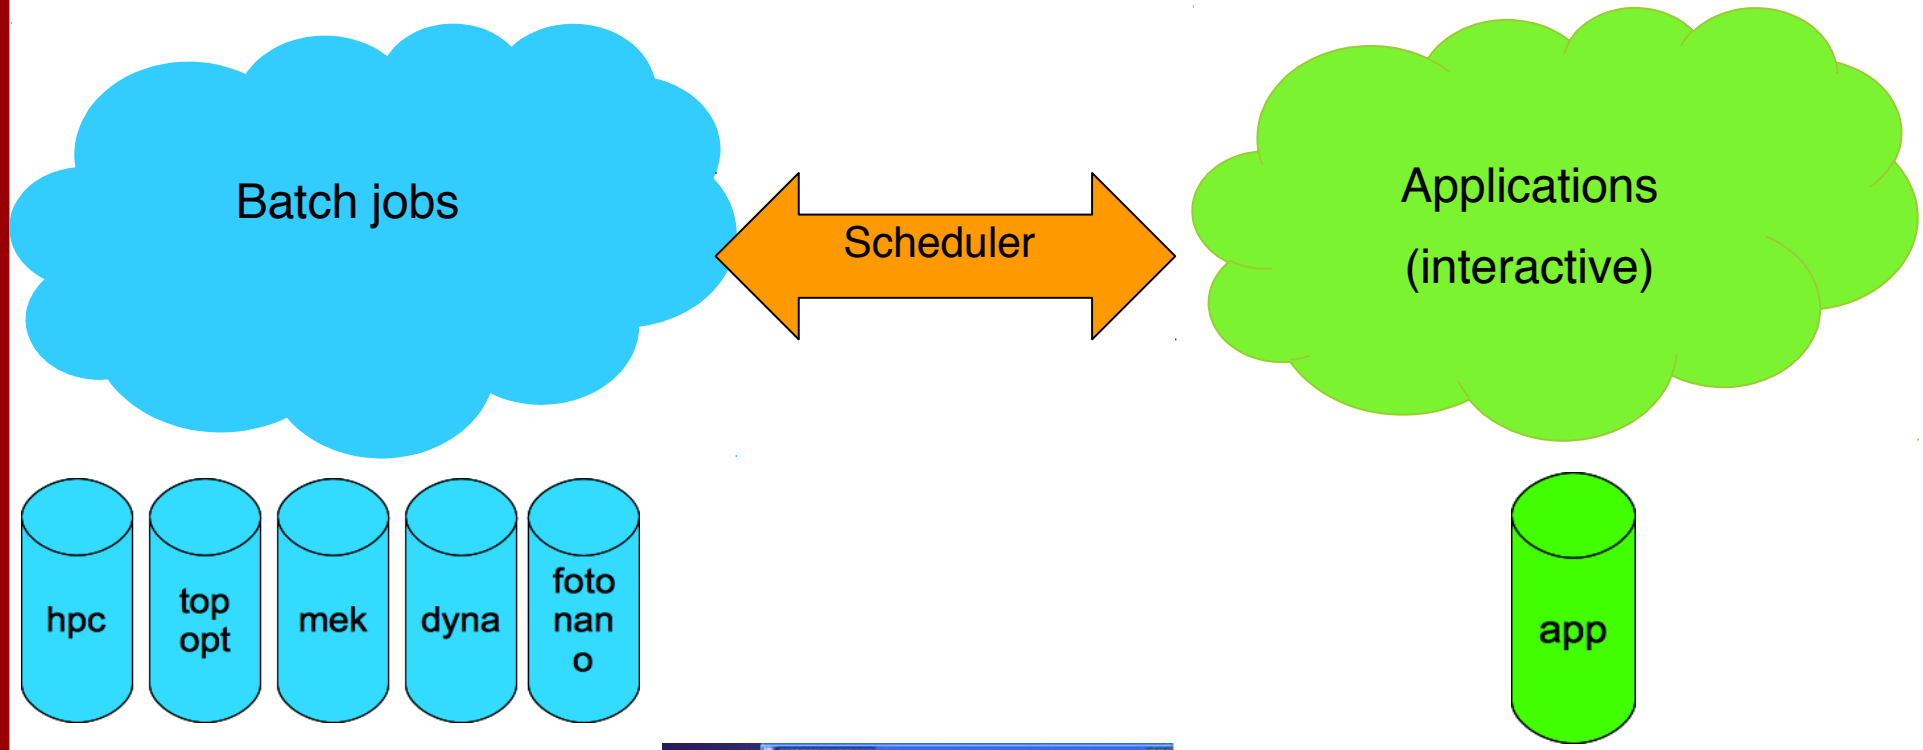
Batch jobs
Scheduler
Applications
(interactive)
hpc
top
opt
mek
dyna
foto
nan
o
app

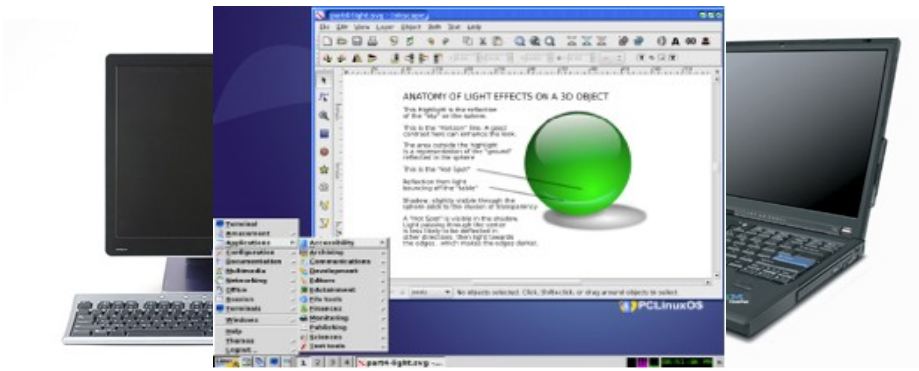
ANATOMY OF LIGHT EFFECTS ON A 3D OBJECT

The central DTU Unix systems

- Application servers – x86_64 based:
  - 6 Huawei XH620 V3 (2x Xeon E5-2660 v3 2.6 GHz)
  - 6 Dell PowerEdge FC430 (2x Xeon E5-2670 v3 2.3 GHz)
  - Scientific Linux 7.x
- Desktop servers (ThinLinc):
  - 3 servers (4x AMD Opteron 6376, 2.4 GHz)
- 10000+ users (students + employees)
  - “everybody has access”!!!

The DTU Unix systems

- HPC servers (for 'everybody'), e.g.
  - 40 IBM NeXtScale nx360 M4 (2x Xeon E5-2680v2 2.8 GHz, 128 GB memory)
  - 15 Huawei XH620 V3 (2x Xeon E5-2660v3 2.6 GHz, 128 GB memory)
  - 40 Huawei XH620 V3 (2x Xeon E5-2650v4 2.2 GHz, 256 GB memory)
  - 24 Lenovo ThinkSystem SD530 (2x Xeon Gold 6126 2.6 GHz, 192-384 GB memory)
- + “private” clusters
  - DTU Compute, DTU Nanotech, DTU Photonics, DTU Chemistry, DTU Elektro, DTU Environment, DTU Management, ...

The DTU Unix systems

- HPC servers DTU Compute
  - 26 HP Proliant SL230s (2x Xeon E5-2665 2.4 GHz, 128 GB memory)
  - 7 IBM NextScale nx360 M5 (2x Xeon E5-2660v3 2.6 GHz, 128 GB memory)
  - 27 Huawei XH620 V3 (2x Xeon E5-2660v3 2.6 GHz, 128 GB memory)
  - 30 Huawei XH620 V3 (2x Xeon E5-2650v4 2.2 GHz, 256 GB memory)
  - 1 node with 1TB memory
  - nodes with NVIDIA V100 GPUs

Interactive Work

Interactive Work

- Starting applications from the desktop menu
- Most desktop apps will run on the same server as your desktop (th1, ..., th3)
- Applications from the DTU menu get dispatched on to one of the application nodes

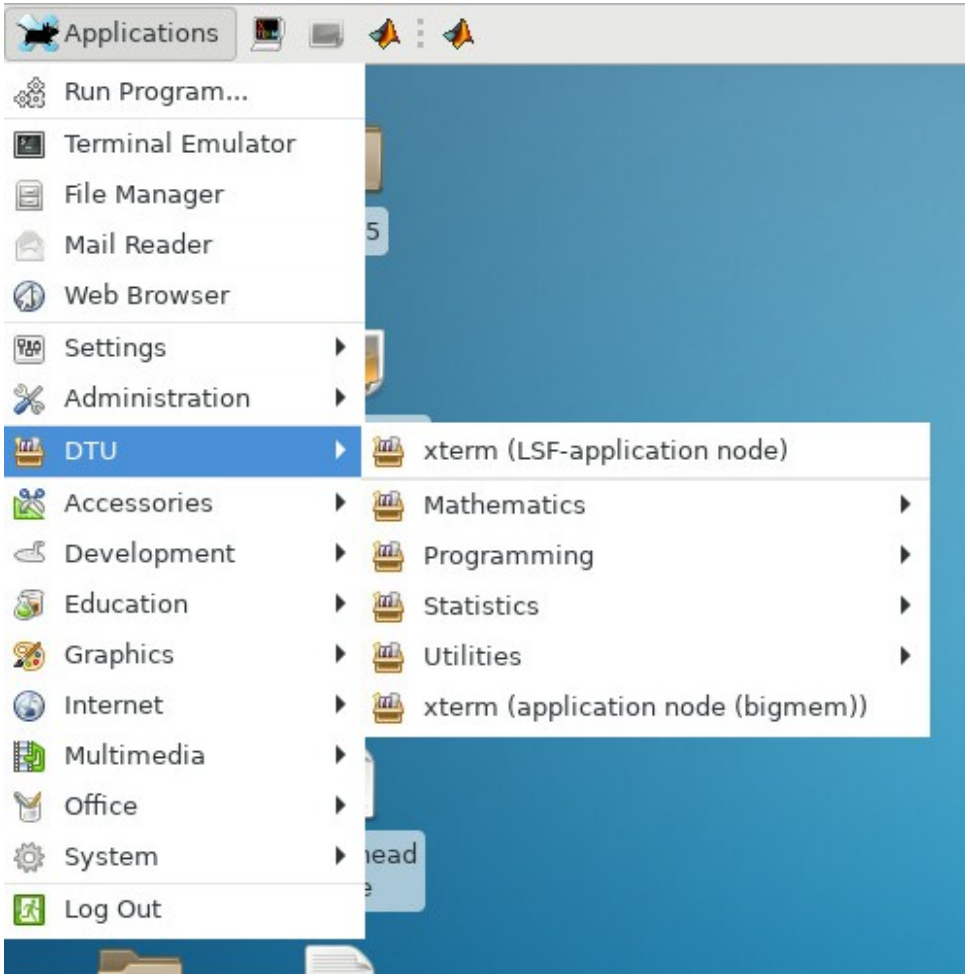

Interactive Work

- “Code of conduct”:
  - applications from the menu run on multi-user systems – so please behave
  - there are resource limits, too:
    - CPU usage: max. 24 hours CPU time
    - memory: max. 16 GB
- If you need more than that, you have to use the batch system!

Access to your files

- To be able to work on the Unix systems, you'll need to transfer your files
- There are several ways to do that:
  - copying: WinSCP, FileZilla, scp, ...
  - mounting your Unix home-directory on your laptop
    - works while on a DTU network, e.g. eduroam
    - for more details see http://gbar.dtu.dk/faq/78-home-directory
- Best practices:
  - avoid spaces and national characters in file names
  - some applications require Unix text format

Software

What is availabe?

Installed Software

- Which software is available?
  - quick answer: a lot!
- How can I find it?
  - can be tricky ... ;-)
  - only the ‘basic’ things are in the desktop menu
  - most programs need to be started from the terminal
  - if in doubt, look in /appl (250+ packages installed)
- Which version of XYZ is installed?
  - check /appl/XYZ for subfolders
  - ... or use ‘modules ...’ (next slide)

Using modules

- modules help to organize certain Unix environment settings, e.g. PATH, MANPATH, LD_LIBRARY_PATH, etc. for different versions of the same application
- list available modules: `module avail`
- load a module: `module load python3`
- swap a version: `module swap python3/3.5.4`
- swap to default: `module swap python3`
- info: http://gbar.dtu.dk/index.php/faq/83-modules

The (new) DCC software stack

- The default module setup is
  - overwhelming
  - confusing
  - not very user friendly

- Try the new DCC software stack:
  - more structured
  - only two releases per year
  - try it: source /dtu/sw/dcc/dcc-sw.bash
  - feedback is welcome!

More Software

- What if it is not there?
  - install in your $HOME folder
  - open a ticket with support, and we will try to help
- Things we do not support:
  - containers (Docker, Singularity) – might come ...
  - Windows software

Lab time ...


THE FUTURE
BELONGS TO
THE FEW OF US
STILL WILLING
TO GET OUR
HANDS DIRTY.

Using the system

- On-line demo:
  - logging in
  - download ZIP file: https://bit.ly/2PbA4T2
    - wget -O ImageXmasWS.zip https://bit.ly/2PbA4T2
  - create folder
  - unzip downloaded file

Batch jobs

Starting point: Take a look at
www.hpc.dtu.dk

Resource Managers

To handle the workload on an HPC installation, one needs a tool to manage and assign the resources: a Resource Manager – sometimes also called 'batch queue system'

- Most common systems:
  - Torque/PBS (ext. scheduler, like Maui or MOAB)
  - LSF
  - Grid Engine
  - Slurm

Resource Managers

Before submitting a job, one has to specify the resources needed, e.g.

- # of CPUs/cores
- amount of memory
- expected run time (wall-clock time)
- other resources, like disk space, GPUs, etc

This is done in a special job script and is system (RM) dependent – but very similar for all RMs.

Resource Managers

- Examples for the DTU batch system, based on
  - ~~MOAB (scheduler) and Torque (resource manager)~~
  - Spectrum LSF (setup used here!)
    - info: HPC User Guides for LSF
- Notes:
  - you need to be on the correct front-end node, either via ThinLinc or ‘ssh login2.hpc.dtu.dk’
  - you cannot submit executables directly, you have to use a job script!
  - don't expect jobs to start immediately – the scheduler has to find free resources first!

Resource Managers

The simplest job script:

submit.sh

```
#!/bin/bash
sleep 60
```

```
$ bsub < submit.sh
Job <702572> is submitted to default queue <hpc>.

$ bstat
JOBID  USER    QUEUE   JOB_NAME SLOTS STAT  START_TIME       ELAPSED
702572 gbarbd  hpc     NONAME       1 RUN   Dec 13 12:17     0:00:00
$ bjobs
JOBID  USER    QUEUE   JOB_NAME SLOTS STAT  START_TIME   TIME_LEFT
702572 gbarbd  hpc     NONAME       1 RUN   Dec 13 12:17 00:15:00 L
$ ls -g
total 4
-rw-r--r-- 1 gbar 1493 Dec 13 12:18 NONAME_702572.out
-rw-r--r-- 1 gbar   22 Dec 13 12:05 simple.sh
```

Resource Managers

The simplest job script – the full story:

simple.sh

```
#!/bin/bash
sleep 60
```

```
$ bsub < simple.sh
bsub info: Job has no name! Setting it to NONAME!
bsub info: Job has no wall-clock time! Setting it to 15 minutes!
bsub info: Job has no output file! Setting it to NONAME_%J.out!
bsub info: Job has no memory requirements! Setting it to 1024 MB!
bsub info:   You need to specify at least -R "rusage[mem=...]"!
Job <702608> is submitted to default queue <hpc>.
```

Resource Managers

A simple job script:

```
#!/bin/bash
#BSUB -J sleeper
#BSUB -o sleeper_%J.out
#BSUB -q hpc
#BSUB -W 2
#BSUB -R "rusage[mem=512MB]"

sleep 60
```

```
$ bsub < submit.sh
Job <702645> is submitted to queue <hpc>.
```

```
$ ls -g
total 3
-rw-r--r-- 1 gbar  121 Dec 13 12:32 submit.sh
-rw-r--r-- 1 gbar 1592 Dec 13 12:36 sleeper_702646.out
```

Resource Managers

- The output file:

```
Sender: LSF System <lsfadmin@n-62-21-20>
Subject: Job 702646: <sleeper> in cluster <dcc> Done

Job <sleeper> was submitted from host <hpclogin3> by user <gbarbd> in
cluster <dcc> at Wed Dec 13 12:34:59 2017.
Job was executed on host(s) <n-62-21-20>, in queue <hpc>, as user
<gbarbd> in cluster <dcc> at Wed Dec 13 12:34:59 2017.
</zhome/.././...> was used as the home directory.
</zhome/.././.../02614/Batch/LSF> was used as the working directory.
Started at Wed Dec 13 12:34:59 2017.
Terminated at Wed Dec 13 12:36:00 2017.
Results reported at Wed Dec 13 12:36:00 2017.

Your job looked like:

------------------------------------------------------------
# LSBATCH: User input
#!/bin/bash
#BSUB -J sleeper
#BSUB -o sleeper_%J.out
```

Resource Managers

- The output file (cont'd):
  - job summary

```
Successfully completed.

Resource usage summary:

    CPU time :                                   0.28 sec.
    Max Memory :                                 4 MB
    Average Memory :                             4.00 MB
    Total Requested Memory :                     512.00 MB
    Delta Memory :                               508.00 MB
    Max Swap :                                   -
    Max Processes :                              4
    Max Threads :                                5
    Run time :                                   65 sec.
    Turnaround time :                            61 sec.

The output (if any) is above this job summary.
```

Resource Managers

Separating output and errors:

```
#!/bin/bash
#BSUB -J sleeper
#BSUB -o sleeper_%J.out
#BSUB -e sleeper_%J.err
#BSUB -q hpc
#BSUB -W 2 -R "rusage[mem=512MB]"

rm nonexistent.txt
echo "Just a minute ..."
sleep 60
```

```
$ bsub < submit2.sh
...
$ ls -g
total 3
-rw-r--r-- 1 gbar  184 Dec 13 13:56 submit2.sh
-rw-r--r-- 1 gbar   63 Dec 13 13:59 sleeper_702793.err
-rw-r--r-- 1 gbar 1744 Dec 13 14:00 sleeper_702793.out
```

Resource Managers

Separating output, errors – and mail summary:

```
#!/bin/bash
#BSUB -J sleeper
#BSUB -o sleeper_%J.out
#BSUB -e sleeper_%J.err
#BSUB -q hpc
#BSUB -W 2 -R "rusage[mem=512MB]"

rm nonexistent.txt
echo "Just a minute ..."
sleep 60
```

send summary at end of job

```
$ bsub -N < submit2.sh
...
$ ls -g
total 3
-rw-r--r-- 1 gbar  184 Dec 13 13:56 submit2.sh
-rw-r--r-- 1 gbar   63 Dec 13 14:04 sleeper_702814.err
-rw-r--r-- 1 gbar   18 Dec 13 14:04 sleeper_702814.out
```

Resource Managers

A simple parallel job script:

- for shared memory (on a single node), using 4 cores

```
#!/bin/bash
#BSUB -J simple_para
#BSUB -o simple_para_%J.out
#BSUB -q hpc
#BSUB -W 2
#BSUB -R "rusage[mem=512MB]"
#BSUB -n 4
#BSUB -R "span[hosts=1]"

export OMP_NUM_THREADS=$LSB_DJOB_NUMPROC
...
```

- Note: the mem=xyzMB is per core!

Resource Managers

more options and examples:

- see http://www.hpc.dtu.dk/ under
  - LSF User Guides
    - http://www.hpc.dtu.dk/?page_id=2534
- do the lab exercises
- use 'man bsub', 'man bjobs', etc

Resource Managers

DTU Computing Center specific commands:

- bstat – shows the status of your jobs; use ‘bstat -h’ for help for other options
- classstat – shows the status of the queues, e.g. free and used cores, pending jobs, etc
- nodestat – shows the current status of all nodes (use 'nodestat hpc' for the nodes of the 'hpc' queue)
- all commands above have a ‘help option’ (-h), but no man-page!

“Big Data”: where to put my (large) datasets?

Large datasets: where & how to

- Limitations of your home folder ($HOME):
  - 30 GB quota
  - ***"slow"*** network file system
- $HOME is backed up
- there are snapshots (hourly, daily, ...) available
- not really suitable for
  - large files
  - temporary files (e.g. during computations)
  - files that change often

Large datasets: where & how to

- The /workN/$USER folders (N = 1,3):
  - a parallel filesystem
  - faster than $HOME
  - only on request
    - but we created one for everybody present today (in /work1)
- Limitations:
  - no backup – and no snapshots(!)
  - good for large files
  - avoid too many files in one folder (true for $HOME, too)
- Don’t use it as permanent storage!!!

Large datasets: where & how to

- How do I transfer my files?
  - use ‘transfer.gbar.dtu.dk’
  - fast uplink from DTU network
  - fast link to all DCC storage systems, i.e. $HOME and /workN
- Do not use transfer host for anything else!
- Do not use the login nodes for large data transfers
  - it works, but it is slow – and it slows down everybody else, too!

Large datasets: where & how to

Other possibilities:

1) access to files that are on some DTU Compute storage
   - possible via project folders under /dtu-compute/
   - needs an agreement with DTU Compute IT
2) shared filesystem(s) with other installations at DTU
   - shared storage with Niflheim (DTU Fysik)
3) bring your own money
   - ... and we'll buy storage for you

Large datasets: where & how to

Please note:

- please don’t upload data that has to follow any GDPR regulations!

Misconceptions & Misunderstandings

The HPC system is always faster!

- Not always true!
- If you have a brand-new laptop, your CPU is probably one or two generations more advanced than the CPUs in the HPC setup ...
- ... but on the HPC system you have access to more resources/cores, than the usual two cores in your laptop!
- Limits: ~100 concurrent cores (batch system)
  - i.e. 100 single core jobs
  - or ten 10-core jobs, or ...

There is more memory available!

- True ... but there are limits as well!
  - we have machines with 64 GB to 256 GB memory
  - Don’t ask for the full memory – leave space for the OS, etc (rule of thumb: ~10%, e.g. 240GB on a 256GB machine)
- Always try to estimate your memory needs, and specify those needs in your job scripts.
- With LSF, we force you to specify your memory needs – and we enforce the limits!
- This helps the scheduler to dispatch to nodes, that can cope with the amount of memory needed.

The more cores – the better ...

- Mostly wrong!
- Your application has to be able to use the cores, i.e. it has to be parallel
  - requesting more than one core for a serial application is a waste of resources
- You might need to activate the parallelism, e.g. via command line options
- Remember Amdahl's law ...

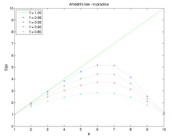

- You'll probably have to run some tests to find the “optimal” number of cores

Amdahl's law

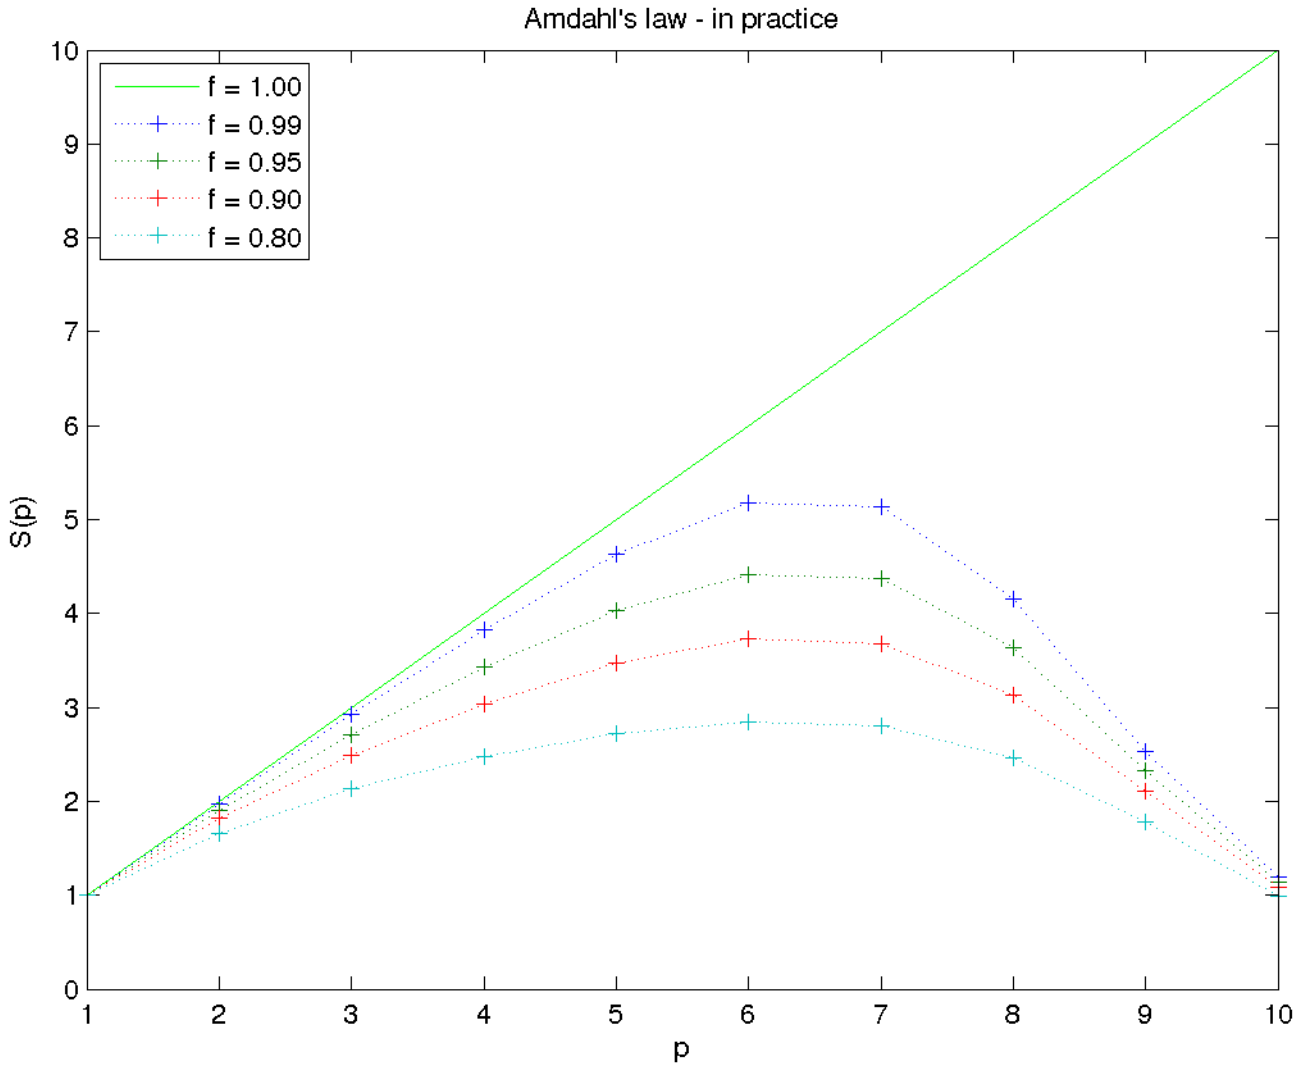
Amdahl's law - in practice
f = 1.00
f = 0.99
f = 0.95
f = 0.90
f = 0.80
S(p)
p
0
1
2
3
4
5
6
7
8
9
10

“Somebody told me that ...”

- ... but it's probably wrong!

- In the support team, we see the same mistakes by users, made over and over again – and that is not productive! Neither for you, nor for the other users, e.g. if you block resources you do not need!
- If you are in doubt, check the web pages, or ask the support people – they are there to help, and just an e-mail away!

“None of my jobs failed ...”

- ... but still there is something wrong!
  - e.g. no output, incomplete output, etc

- LSF reports job success (DONE) or failure (EXIT)
- this information is taken from the return code of the last command in your batch script!
- Things like “echo done.” or “date” will always return success – and that will hide problems!
- Conclusion: make your application the last command in the batch script!

Summary

Summary – DTU Unix Systems

- a general computing resource
- Linux based
- accessible
  - from everywhere
  - for everybody with a DTU account
- extends the compute power of your own computer/laptop
- Not to forget: your data is backed up!

Where to get help?

- Today:
  - ask us – we are here to help!
- Later:
  - check our webpages:
    - www.hpc.dtu.dk
    - www.gbar.dtu.dk
  - send e-mail to:
    - support@hpc.dtu.dk (HPC related help)
    - support@cc.dtu.dk (general help)